

INSIDE
TODAY

◆ SECURITY

Three Kubernetes CVEs in a week, one hits multi-tenancy P.2

◆ HPC

Azimuth adds Kubernetes 1.36 and AMD GPUs, drops Ingress NGINX P.2

◆ OPENSTACK

OpenStack 2026.2 "Hibiscus" ships with 28% more changes P.2

KUBERNETES

Kubernetes 1.37 lands 67 enhancements, from HPA scale-to-zero to stable KYAML

Sixteen features go stable, 23 move to beta and 27 debut in alpha. These are the ones that change day-to-day work for cluster and platform operators.

COMPILED FROM KUBERNETES.IO, THE NEW STACK

The release ships 67 enhancements: 16 stable, 23 beta, 27 alpha and one deprecation. The headline stable feature is KYAML, a YAML dialect designed for manifests that removes two long-standing traps:


whitespace sensitivity and implicit typing of values.

HorizontalPodAutoscaler scale-to-zero reaches beta and is on by default for workloads scaled on object or external metrics, so idle pods can drop to zero. Manifest-based admis-

sion control also moves to beta, Pod-level checkpoint and restore debuts in alpha, and resilient watch cache initialisation is complete, to prevent control-plane outages in large clusters.

Continued on page 2

1.37 BY THE NUMBERS



67 enhancements in total. Source: The New Stack.

AI ON KUBERNETES

One control plane for 10,000 accelerators: how China Merchants Bank runs AI on Kubernetes

Training, fine-tuning and inference now share the same hardware, and inference costs fell by more than 60%.

COMPILED FROM CNCF

The winner of CNCF's end-user case study contest pools nearly 10,000 heterogeneous accelerator cards, 99% of the bank's accelerator resources, under one cloud native framework. Training, fine-tuning and on-line inference keep separate runtimes but share the hardware.

Each project has one job. Kueue handles admission, queues and quotas, so training jobs cannot reserve capacity before they run. KEDA scales inference on live demand. HAMI splits accelerators into fine-grained units for both paths, Fluid speeds up access to datasets and model weights, and Prometheus covers observability.

Average accelerator utilisation rose from 35% to over 60%, and inference cost per million tokens fell by more than 60%. The bank's Twinkle framework lets five LoRA tenants share one base model, cutting accelerator use by 80% and raising training density fivefold.

Full story → www.cncf.io/announcements/2026/09/07/china-merchants-bank-wins-cncf-end-user-case-study-contest-for-unifying-ai-training-and-inference-on-kubernetes

THE BANK IN NUMBERS

~10k accelerator cards on one control plane
60%+ average utilisation, up from 35%
-60% inference cost per million tokens

PLATFORM ENGINEERING

Yassir runs access to 15 tools with Crossplane and Backstage


COMPILED FROM MEDIUM • AYOUB REDA

Under Terraform, reviews never checked whether a change made organisational sense, and a single state file queued even unrelated changes. Now there are two CRDs, Member and Team, plus one Crossplane composition per tool: native providers for GCP, GitHub, PagerDuty and Atlas, the HTTP provider for Slack, Jira and Sentry.

Continuous reconciliation strips any access granted by hand. CODEOWNERS is generated from team files, privileged grants are blocked in CI, and Backstage forms open the same pull requests an engineer would. Compositions are tested with crossplane render in CI. The authors' lesson: getting the right approvers mattered more than the tool.

Full story → medium.com/p/0c744f10b671

CARTOON OF THE DAY



Drawing: Claude

Water can wait. First, the password.

WHY IT'S FUNNY An exhausted traveller crosses the desert and finally reaches the oasis. He ignores the pond under the palm tree and crawls towards the free Wi-Fi sign instead. The joke is about how connectivity has become a "basic need", almost ahead of water.

DATES TO WATCH

Events and calls for papers

Oct 8
Dapr Day Virtual

ONLINE

Oct 11
KubeCon + CloudNativeCon Europe 2027 (Barcelona, Mar 15-18): talks due 23:59 CEST

CFP

Oct 15
KubeVirt Summit Virtual

ONLINE

NEWSROOM

PRESS RELEASE · FROM GITHUB RELEASES AND NOTIFICATIONS

Blueprint rendering gets a new name and offline previews; Autopilot isolates sub-agent sessions

October 2, 2026 — Two render-service releases on the public side; one merged change on the private side.

EXAMPLE PLATFORMOPS

GitHub releases, github.com/example-org

helm-render-service 0.5.0 · Sep 30

- Chart renamed to blueprint-render-service, new OCI registry.
- Endpoint secret default is now blueprint-render-endpoint: update values before upgrading.

github.com/example-org/<repo>/releases

helm-render-service 0.4.0 · Sep 28

- lookupStubs on /render: previews without a live cluster. Pre-flight checks in CI.

github.com/example-org/<repo>/releases

Latest elsewhere: component-a 1.12.13 (Sep 24), component-b 0.4.0 (Sep 22)

EXAMPLE PRIVATE SUITE

Merged into main this week

autopilot PR #94 · merged Oct 2

- Each sub-agent call gets its own session; isolateSessions is live on kagent 0.10.1.

No other private pull requests were merged in the last seven days.

CONTINUED FROM PAGE 1 · KUBERNETES

Kubernetes 1.37

On nodes, Memory QoS graduates to beta and needs cgroup v2. Pod-level resource managers reach beta but stay off by default: they read the Pod's .spec.resources, so lightweight sidecars no longer need exclusive cores to get NUMA alignment. Tracking when a PVC was last used goes beta, which helps find forgotten volumes, and new options harden emptyDir permissions and bind mounts.

Full story → thenewstack.io/kubecon-kubernetes-updates-security
kubernetes.io/blog

OPENSTACK

OpenStack 2026.2 “Hibiscus” ships with 28% more changes

The 34th release, out on September 30, is built around AI-era infrastructure.

COMPILED FROM OPENSTACK BLOG

About 600 contributors landed some 11,500 changes, up 21% and 28% on the Gazpacho cycle. Neutron cut the memory footprint of HA router monitoring by more than fifteen times. The cycle's themes are GPU and accelerator support and confidential computing, and the Vulnerability Management Team coordinated dozens of embargoed advisories.

Hibiscus is not a SLURP release. The next one, Indri, is, so operators on Gazpacho (2026.1)

can skip Hibiscus and upgrade straight to Indri.

Full story → legacy.openstack.org/blog/openstack-hibiscus-what-it-takes-to-build-a-release-together

NEW MODELS

GPT-6.1 Sol, Claude Sonnet 5.5 and open-weight MiMo-V2.6 land in one week

Frontier performance keeps getting cheaper.

COMPILED FROM THE NEW STACK, UNITE.AI, DATANORTH

OpenAI's GPT-6.1 Sol (September 29) costs \$2 per million input tokens and \$10 per million output, a fifth of GPT-6 Astra, and OpenAI says it nearly matches Astra on agentic coding and computer use. It is in the API and ChatGPT's paid tiers.

Anthropic's Claude Sonnet 5.5 (September 28) keeps Sonnet 5 pricing at \$2 and \$10. Anthropic says it is over 30% faster and up to 30% cheaper per task, with adaptive thinking on by default; it runs on AWS, Google Cloud and Azure.

Xiaomi's MiMo-V2.6 (September 21) is the open option: MIT-licensed weights on Hugging Face, a one-million-token context, a Pro model with 1.02 trillion parameters (42 billion active) and a 309-billion-parameter Flash. Xiaomi gives no hardware guidance for self-hosting.

Full story → thenewstack.io/openai-gpt-6-1-sol

www.unite.ai/anthropic-releases-claude-sonnet-5-5-at-unchanged-sonnet-5-pricing
datanorth.ai/news/xiaomi-releases-mimo-v2-6-pro-and-flash

NEWS IN BRIEF

SECURITY

Three Kubernetes CVEs in a week, one hits multi-tenancy

September 23: CVE-2026-2270, write access to StatefulSets and ControllerRevisions allows creating pods in other namespaces; CVE-2026-76654, sub-Path symlinks on Windows nodes allow NTLM coercion. September 28: CVE-2026-19444, path traversal in kubectcl cp on Windows enables arbitrary file writes. Check the advisories for severity and fixed versions.

— [kubernetes-security-announce groups.google.com/g/kubernetes-security-announce](https://kubernetes-security-announce.groups.google.com/g/kubernetes-security-announce)

HPC

Azimuth adds Kubernetes 1.36 and AMD GPUs, drops Ingress NGINX

StackHPC's Q3 report (September 7): Azimuth 2026.7 brings Kubernetes 1.36 and AMD GPU support; 2026.6 replaced the archived Ingress NGINX with Traefik, with some downtime; 2026.8 moves to Ubuntu 24.04, adds single-node control planes and swaps the deprecated Kubernetes dashboard for Headlamp. Five StackHPC customer systems are on the TOP500, including Sunrise (#67) and Dawn (#111).

— StackHPC
www.stackhpc.com/navigating-upstream-2026-Q3.html


PUZZLE · WORD SEARCH

Find the words from your day job hidden in the grid.

FIND CLUSTER · CROSSPLANE · GITOPS · HELM · KARMADA · KUBERNETES · KUEUE · OPENSTACK · POD · SLURM

Y	W	H	C	R	T	T	V	P	W	M	L
I	R	O	R	E	M	I	T	A	S	I	S
L	U	P	O	T	K	R	L	A	W	N	A
N	O	E	S	S	U	O	U	R	P	S	N
N	F	N	S	U	B	A	W	L	A	N	S
C	V	S	P	L	E	K	M	H	S	P	Y
G	M	T	L	C	R	A	A	O	T	Y	N
I	L	A	A	N	N	R	K	U	E	U	E
T	E	C	N	D	E	M	B	V	T	D	U
O	H	K	E	O	T	A	R	A	D	I	W
P	V	I	S	M	E	D	C	N	F	O	M
S	S	U	A	M	S	A	N	S	F	E	P

Words run in every direction, also backwards.



NOUJNTO5