



INSIDE TODAY

◆ KUBERNETES

Patch window closes October 9

P.2

◆ CLOUD NATIVE SECURITY

A free post-quantum CA, in Merkle trees P.2

◆ NETWORKING

Cilium lets two pods share one LB IP P.2

AI ON KUBERNETES

Local SSD swap triples agent sandbox density: 80 pods a node become 240

Node swap has been GA since Kubernetes v1.34 and now has numbers behind it. The win is real, but it has a floor: squeeze too hard and run times jump 40%.

COMPILED FROM KUBERNETES.IO

A Kubernetes blog post of October 5, by Ocean Xie and Yuan Wang, finally measures node swap, GA since v1.34. On one c4-standard-32 node, 32 vCPU and 120 GB of RAM, with swap on Local SSD, Python agent

sandboxes went from 80 concurrent sessions to 240.

The gain is idle memory, not faster hardware: agents allocate a large footprint, then wait for a prompt. gVisor sandboxes doubled, 80 pods to 160; runc went 512 to 768; Kata Contain-

ers, with their own guest kernel, only moved 40 to 50.

There is a floor. A kernel build capped at 300 MB with swap beat its 600 MB baseline, 374 seconds against 433, but at 200 MB it lost 40% to I/O wait. The kubelet wants failSwapOn:

Continued on page 2

PODS PER NODE, WITH AND WITHOUT SWAP

80 GVISOR, NO SWAP

160 GVISOR + SWAP

512 RUNC, NO SWAP

768 RUNC + SWAP

Python agent sandboxes went from 80 to 240 concurrent sessions on one c4-standard-32, 32 vCPU and 120 GB of RAM. Source: kubernetes.io.

OPENSTACK

OpenStack 2026.2 Hibiscus arrives with Intel TDX, AMD SEV-SNP and Lustre in Manila

The 34th release closed on September 30. It is not a SLURP, so operators can sit on it and skip straight to 2027.1 Indri in March.

COMPILED FROM OPENINFRA FOUNDATION

Hibiscus shipped on September 30 after a six-month cycle with about 600 contributors, 21% more than Gazpacho, and 11,500 code changes, up 28%. Nova extends confidential computing to AMD SEV-SNP and Intel TDX, and with Cyborg manages mediated devices: vGPU slices handed to tenants.

The HPC-facing change is Manila, which now speaks to Weka and Lustre parallel filesystems. Neutron cut the memory footprint of HA router monitoring by more than 15x, Glance gained parallel image imports, and Designate adds post-quantum tooling.

Full story → www.openstack.org/software/openstack-hibiscus/releases.openstack.org/hibiscus

AI INFERENCE

vLLM 0.31.0 rejects per-request multimodal kwargs and drops the slow tokenizer

717 commits from 307 contributors, released October 5. Two changes will break serving deployments on upgrade.

COMPILED FROM GITHUB • VLLM

FlashMLA mega attention with the DeepSeek V4.1 NVFP4 compressed KV cache is now the SM100 default. A vllm preload CLI keeps quantized weights in GPU memory across engine restarts, and

--max-num-active-seqs caps

Continued on page 2

OBSERVABILITY

Kubelet only watches inodes at 5% free, and by then it is evicting pods


A CNCF mentorship post of October 5 traces a node that burned 11 million inodes on container snapshots.

COMPILED FROM CNCF

Disk bytes get background image garbage collection at 85% full. Inodes get nothing until nodefs.inodesFree hits its 5% hard threshold, 95% usage, and then pods are evicted. The node described sat at 67% inode usage against 83% disk, 11M of roughly 16.3M, because one Node.js image held 21,553 files in a single node_modules package and each rebuild made a new snapshot. Alert on a flat 80% inode threshold.

Full story → www.cncf.io/blog/2026/10/05/kubelet-watches-inodes-just-not-until-its-an-emergency

CARTOON OF THE DAY



Drawing: Claude

Et tu, alarm clock?

WHY IT'S FUNNY An old-fashioned bell alarm clock has grown a pair of arms and is reaching up to press its own snooze button, with Zs drifting away from it. The joke is that the one object in the house whose entire purpose is to wake somebody up has decided it would rather have another ten minutes. Everyone who has ever hit snooze recognises the betrayal.

DATES TO WATCH

Events and calls for papers

Oct 7-8

PromCon Europe, Munich

EVENT

Oct 8

Dapr Day Virtual

ONLINE

Oct 9

NextGen Intelligence: Data, Cloud & AI (London, Nov 17): talks due

CFP

PRESS RELEASE · FROM GITHUB RELEASES AND NOTIFICATIONS

X O B D N A S S
 M A T E R M A R K U
 A S Q
 L E U Y C
 I D A T I S E T
 N O N I B E L
 V A N
 M I N T S
 W I U M S
 M A N O R T U E N
 d